

Introduction to Mahout

Yun Liaw

Vpon Inc.

What is Mahout

A person who drives an elephant(?)



Mahout is...

- A machine learning libraries
- Scalable and robust
 - Scalable to large data
 - Scalable to support business cases
- Based on Hadoop
 - MapReduce Based
 - But will shift to Apache Spark

Mahout's Currently Use Cases

- Classification
- Clustering
- Recommendation

Classification Algorithms

- Naive Bayes / Complementary Naive Bayes
- Random Forest
- Logistic Regression (trained via SGD)
 - Online learning
- Hidden Markov Models
- Multilayer Perceptron

Clustering Algorithms

- Canopy Clustering
 - deprecated, will be removed once Streaming k-Means is stable enough
- L-Means Clustering
- Fuzzy k-Means
- Streaming k-Means
- Spectral Clustering

Recommendation Algorithms

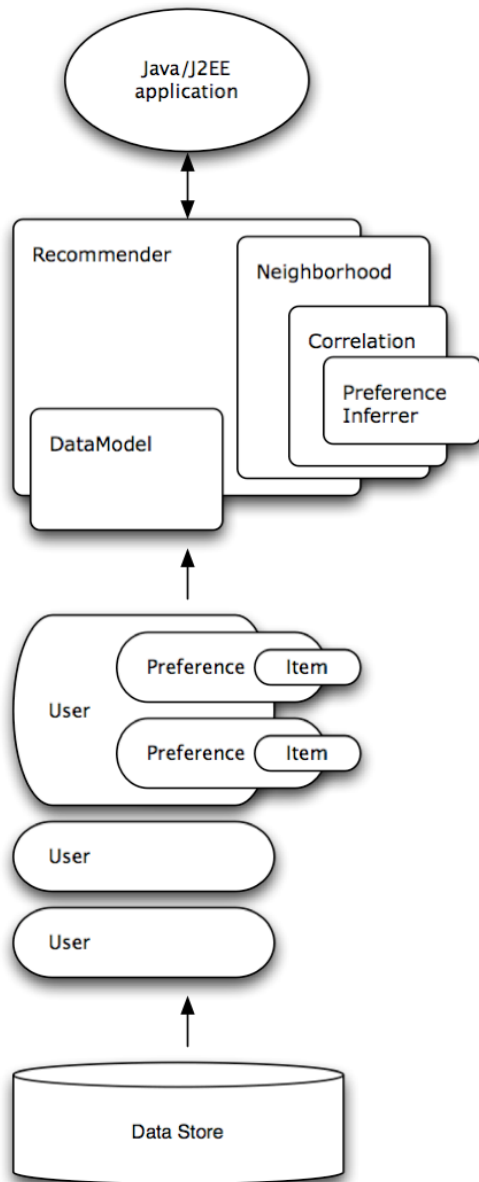
- User-Based Collaborative Filtering
- Item-Based Collaborative Filtering
- Matrix Factorization with Alternating Least Squares (Implicit Feedback)
- Weighted Matrix Factorization, SVD++, Parallel SGD

Others...

- Dimensionality Reduction
 - Singular Value Decomposition
 - Lanczos Algorithm
 - Stochastic SVD
 - Principal Component Analysis (via Stochastic SVD)
- Topic Models
 - Latent Dirichlet Allocation
- Miscellaneous
 - Frequent Pattern Mining
 - RowSimilarityJob
 - compute pairwise similarities between the rows of a matrix
 - ConcatMatrices
 - combine 2 matrices or vectors into a single matrix
 - Collocations
 - find co-locations of tokens in text

Recommender!

Recommendation Architecture



A Java/J2EE Application that invokes..

Mahout Recommendation Engine..

Whose DataModel is based on a set of user preferences..

That are built on the ground of physical datasource

Recommendation in Mahout

- Input: raw data (user preference)
- Output: preferences estimation
- Step1: raw data -> DataModel
- Step2: Recommendation
- Step3: Compute rating estimation
- Step4: Evaluating Recommendation



Recommendation Components


- DataModel Interface
 - Methods for mapping raw data to a Mahout-compliant form
- Similarity Interface
 - Methods to calculate the degree of correlation between two objects (user, item, etc.)
- Neighborhood Interface
 - Methods to define the concept of ‘neighborhood’
- Recommender Interface
 - Methods to implement the recommendation step itself

Component: DataModel

- Example: FileDataModel

Consider a CSV File:

1, 101, 5.0

1, 102, 3.0  User 1 has preference 3.0 for item 102

2, 101, 2.0

2, 102, 2.5

2, 103, 3.0

2, 104, 5.0

...

Basic object: Preference

A Triple Object -> (user, item, rating)

Stored in **UserPreferenceArray**


Component: Similarity

- Many different similarity measures
 - Pearson Correlation
 - Euclidean Distance
 - LogLikelyhood
 - etc.
- Different Similarity influence the **Neighborhood** formation


Component: Neighborhood

- Two definitions of Neighborhood:
 - Nearest N users
 - The first N users with the highest similarity are labeled as 'neighbors'
 - Thresholds
 - Users whose similarity is above a threshold are labeled as 'neighbors'

Nearest N Users



Threshold



Component: Recommender

- Given a **DataModel**, a **Similarity** measure between objects (user or item), and a definition of **Neighborhood**, a **recommender** can produce the preference estimation as the output
- Recommenders:
 - User-based CF
 - Item-based CF
 - SVD (single value decomposition)
 - ALS (alternating least square)
 - etc.

Evaluation

- Prediction-based estimator
 - (MAE) Mean Absolute Error
 - RMSE (Root Mean Square Error)
- IR (information retrieval) based estimator
 - Precision = $TP / (TP + FP)$
 - Recall = $TP / (TP + FN)$
 - F1Measure = $2PR / (P + R)$

Let's Try Mahout!

Environment Setup

- Download Eclipse
 - <http://www.eclipse.org/downloads/>
- Install Maven plugin for Eclipse
 - <http://www.eclipse.org/m2e/download/>
- Create Maven Project in Eclipse
- Add Maven Dependency
 - <https://mahout.apache.org/users/basics/quickstart.html>
- HADOOP IS NOT MANDATORY

JAVADOC

<https://builds.apache.org/job/Mahout-Quality/javadoc/>

Exercise 1

- Create a Preference object
- Set some preferences
- Print it!
- Hint:
 - GenericUserPreferenceArray
 - Preference

Exercise 2

- Create DataModel that stores multiple user's preference
- PreferenceArray stores preference of a single user
 - HashMap?
- Hint:
 - FastByIDMap

Exercise3

- Create DataModel
- Feed the data from CSV file
- Calculate Pearson Cooreation and Euclidean Distance
- Hint:
 - FileDataModel
 - PearsonCorrelationSimilarity
 - EuclideanDistanceSimilarity

Exercise4

- Create a DataModel
- Feed the DataModel through a CSV file
- Generate Neighborhood
- Generate Recommendations
- Hint:
 - NearestNUserNeighborhood
 - GenericUserBasedRecommender

Exercise5

- Let's start with some real data!
 - Download GroupLens Dataset
 - <http://files.grouplens.org/datasets/movielens/ml-100k.zip>
- Play this dataset with Exercise3 and Exercise4!
- Tweak the parameters
 - Different similarity measures
 - Different Neighborhood
- But which one is better?
 - We need evaluation!

Exercise6

- Evaluate different recommender configurations on MovieLens data
- Metrics: MAE, RMSE
- Hint:
 - RandomUtils.useTestSeed()
 - To ensure that random seed is the same
 - RecommenderEvaluator
 - AverageAbsoluteDifferenceRecommenderEvaluator (MAE)
 - RMSRecommenderEvaluator (RMSE)
 - RecommenderBuilder
 - To generate the recommender instance

Exercise7

- Evaluate different recommender configurations on MovieLens data
- Metrics: Precision, Recall, F1-Score
- Hint:
 - GenericRecommenderIRStatsEvaluator

Exercise8

- Item-based Recommender
- Hint:
 - ItemSimilarity
 - GenericItemBasedRecommender

Running on Hadoop

- Package into Jar
 - Maven will help A LOT
- Using Hadoop jar command to run

```
hadoop jar your_jar.jar \  
org.apache.mahout.cf.taste.hadoop.pseudo.RecommenderJob \  
-Dmapred.input.dir=hdfs/position/for/input/data \  
-Dmapred.output.dir=output \  
--recommenderClassName \  
your.own.implement.very.special.AwsomeRecommender
```

Future

- Hadoop -> Spark
 - <http://spark.apache.org/>
 - Less code with Scala implementation
 - Faster computation
 - Utilize memories in cluster
 - Machine learning is built in core (MLLib)
 - With streaming processing (Spark Streaming)
 - With Hive-compatible SQL (Spark SQL)
 - With graph algorithms (Graph X)
 - AWSOME!

Mahout Resources

- Official Site
 - <https://mahout.apache.org/>
- JAVADOC
 - <https://builds.apache.org/job/Mahout-Quality/javadoc/>
- Github
 - <https://github.com/apache/mahout>
- Books
 - [Mahout in Action \(2011\)](#)
 - [Apache Mahout Cookbook \(2013\)](#)